

Statistical Learning for Data Science

MSDS534

Lecture 05: Kernel Methods

Department of Statistics & Biostatistics
Rutgers University

Acknowledgement

- Some of the figures in this presentation are taken from *An Introduction to Statistical Learning, with applications in R* (Springer, 2013) with permission from the authors: G. James, D. Witten, T. Hastie and R. Tibshirani.
- Some of the figures in this presentation are taken from *Elements of Statistical Learning* (Springer, 2009) with permission from the authors: T. Hastie, R. Tibshirani and J. Friedman.
- Some of the figures in this presentation are taken from *Pattern Recognition and Machine Learning* (Springer, 2006) with permission from the author: Christopher M. Bishop.
- Some of the Lab codes in this presentation are taken from *An Introduction to Statistical Learning, with applications in R* (Springer, 2013) with permission from the authors: G. James, D. Witten, T. Hastie and R. Tibshirani.

Reading Assignments

- ESL: § 14.5.1, § 14.5.3, § 14.5.4
- ISL: § 10.2, § 10.4, § 10.5
- PRML: § 12.3

1 Review: Principal Component Analysis

2 Kernel PCA

3 Spectral Clustering

Principal Component Analysis

- Suppose we have N training points $\mathbf{x}_1, \dots, \mathbf{x}_N \in \mathbb{R}^p$. Denote by $\mathbf{X} = (\mathbf{x}_1, \dots, \mathbf{x}_N)'$ the $N \times p$ data matrix.
- Find a direction along which the data has the largest variation.

$$\max_{\|\mathbf{v}\|=1} \{\text{sample variance of } \mathbf{X}\mathbf{v}\}.$$


- Find a line that is “closest” to the data.
 - A line in \mathbb{R}^p can be represented as $f(\eta) = \boldsymbol{\mu} + \eta\mathbf{v}$, where $\boldsymbol{\mu}$ and \mathbf{v} are p -dimensional parameter vectors.
 - For a point \mathbf{x}_i , we find η_i such that

$$\eta_i = \arg \min_{\eta} \|\mathbf{x}_i - (\boldsymbol{\mu} + \eta\mathbf{v})\|^2.$$



- Next, find the “closest” line:

$$\min_{\boldsymbol{\mu}, \|\mathbf{v}\|=1} \sum_{i=1}^N \|\mathbf{x}_i - (\boldsymbol{\mu} + \eta_i\mathbf{v})\|^2.$$

Principal Component Direction



Best Linear Approximation



Principal Component Analysis

Orthogonal Matrix

Let \mathbf{A} be a $n \times m$ matrix with $n \geq m$. Denote its m columns by $\mathbf{a}_1, \dots, \mathbf{a}_m$. We say \mathbf{A} is **orthogonal** if its **columns** are **orthonormal**, i.e.

$$\mathbf{a}_i' \mathbf{a}_j = \begin{cases} 1 & \text{if } i = j; \\ 0 & \text{if } i \neq j. \end{cases}$$

- Having find $\hat{\mathbf{v}}_1$, find next one, **subject to the constraint that it is orthogonal to $\hat{\mathbf{v}}_1$** .
- Find the best low dimensional linear approximations to the data. Consider the rank- q linear model for representing the p -dimensional data $\mathbf{x}_1, \dots, \mathbf{x}_N$.

$$f(\boldsymbol{\eta}) = \boldsymbol{\mu} + \mathbf{V}_q \boldsymbol{\eta},$$

where $\boldsymbol{\mu} \in \mathbb{R}^p$ is a location vector, \mathbf{V}_q is a $p \times q$ orthogonal matrix, and $\boldsymbol{\eta} \in \mathbb{R}^q$ is a vector of parameters. Fitting such a model to the data by least squares amounts to minimizing the **reconstruction error**

$$\min_{\boldsymbol{\mu}, \{\boldsymbol{\eta}_i\}, \mathbf{V}_q} \sum_{i=1}^N \|\mathbf{x}_i - \boldsymbol{\mu} - \mathbf{V}_q \boldsymbol{\eta}_i\|^2.$$

PCA: Preprocessing

- 1 Compute $\bar{\mathbf{x}} = \frac{1}{N} \sum_{i=1}^N \mathbf{x}_i$, and then subtract the sample mean from each observation

$$\mathbf{x}_i := \mathbf{x}_i - \bar{\mathbf{x}}.$$

- 2 Optional. Preferred when features are on different scales. Normalize each feature.

$$\hat{s}_j = \sqrt{\sum_{i=1}^N x_{ij}^2} \quad \text{then} \quad x_{ij} := \frac{x_{ij}}{\hat{s}_j} \quad \text{for all } 1 \leq i \leq N, 1 \leq j \leq p.$$

From now on we always assume Step 1 has been done. Problems become

- Direction of maximum sample variance

$$\max_{\|\mathbf{v}_1\|=1} \mathbf{v}_1' \mathbf{X}' \mathbf{X} \mathbf{v}_1.$$

- Best linear approximation

$$\min_{\{\eta_i\}, \mathbf{v}_1} \sum_{i=1}^N \|\mathbf{x}_i - \mathbf{v}_1 \eta_i\|^2.$$

Principal Directions & Best Linear Approximation

- Principal directions. Having found $\mathbf{v}_1, \dots, \mathbf{v}_j$, find \mathbf{v}_j by

$$\begin{aligned} \max_{\mathbf{v}} \quad & \mathbf{v}' \mathbf{X}' \mathbf{X} \mathbf{v} \\ \text{subject to: } & \mathbf{v} \perp \mathbf{v}_1, \dots, \mathbf{v}_{j-1}. \end{aligned}$$

- Best linear approximation. Find the optimal orthogonal matrix $\mathbf{V}_q \in \mathbb{R}^{p \times q}$,

$$\min_{\{\boldsymbol{\eta}_i\}, \mathbf{V}_q} \sum_{i=1}^N \|\mathbf{x}_i - \mathbf{V}_q \boldsymbol{\eta}_i\|^2.$$

- Fix \mathbf{V}_q , we must have

$$\hat{\boldsymbol{\eta}}_i = \mathbf{V}_q' \mathbf{x}_i.$$

- The problem is reduced to

$$\min_{\mathbf{V}_q} \sum_{i=1}^N \|\mathbf{x}_i - \mathbf{V}_q \mathbf{V}_q' \mathbf{x}_i\|^2.$$

Singular Value Decomposition

The **singular value decomposition (SVD)** of the $N \times p$ (**assume $N \geq p$**) matrix \mathbf{X} has the form

$$\mathbf{X} = \mathbf{U}\mathbf{D}\mathbf{V}'.$$

- \mathbf{U} and \mathbf{V} are $N \times p$ and $p \times p$ orthogonal matrices.
- The columns of \mathbf{V} , denoted by $\mathbf{v}_1, \dots, \mathbf{v}_p$, span the row space of \mathbf{X} .
- The columns of \mathbf{U} , $\mathbf{u}_1, \dots, \mathbf{u}_p$, span the column space of \mathbf{X} .
- \mathbf{D} is a $p \times p$ diagonal matrix, with diagonal entries $d_1 \geq d_2 \geq \dots \geq d_p \geq 0$, which are called **singular values** of \mathbf{X} .
- $\mathbf{v}_1, \dots, \mathbf{v}_p$ are called **right singular vectors**. They are also eigenvectors of the matrix $\mathbf{X}'\mathbf{X}$, corresponding to the eigenvalues $d_1^2 \geq d_2^2 \geq \dots \geq d_p^2$.
- $\mathbf{u}_1, \dots, \mathbf{u}_p$ are called **left singular vectors**. They are also eigenvectors of the matrix $\mathbf{X}\mathbf{X}'$, corresponding to the eigenvalues $d_1^2 \geq d_2^2 \geq \dots \geq d_p^2$. The rest eigenvalues of $\mathbf{X}\mathbf{X}'$ are all zero.
- It can be rewritten as $\mathbf{X} = d_1\mathbf{u}_1\mathbf{v}_1' + \dots + d_p\mathbf{u}_p\mathbf{v}_p'$.
- Such a decomposition is unique (up to a sign change) if all d_j are distinct.

- Compute the singular value decomposition (SVD) of $\mathbf{X} = (\mathbf{x}_1, \dots, \mathbf{x}_N)'$

$$\mathbf{X} = \mathbf{U}\mathbf{D}\mathbf{V}'.$$

- The principal directions are the right singular vectors $\mathbf{v}_1, \dots, \mathbf{v}_p$.
- For the best linear approximation problem, for each $1 \leq q \leq p$, the optimal $\mathbf{V}_q = (\mathbf{v}_1, \dots, \mathbf{v}_q)$.

 \mathbf{v}_m m -th principal direction $\mathbf{z}_m = \mathbf{X}\mathbf{v}_m = d_m\mathbf{u}_m$ m -th principal component \mathbf{v}_m loadings of the m -th principal component

*Sparse PCA

- Joliffe et al (2003). [SCoTLASS](#).

$$\max v' \mathbf{X}' \mathbf{X} v \quad \text{subject to } \|v\| = 1, \sum_{j=1}^p |v_j| \leq t.$$

- Not convex means the computations are difficult.


- Zou et al (2006). [Sparse PCA](#).

$$\min_{\theta, \mathbf{v}_1} \sum_{i=1}^N \|\mathbf{x}_i - \theta \mathbf{v}_1' \mathbf{x}_i\|^2 + \lambda \|\mathbf{v}_1\|_2^2 + \lambda_1 \|\mathbf{v}_1\|_1 \quad \text{subject to } \|\theta\|_2 = 1.$$



- Fix θ , minimization over \mathbf{v}_1 is equivalent to [elastic-net](#) problems, which can be solved by LARS.
- Fix \mathbf{v}_1 , minimization over θ is done by a SVD calculation.

Example

- Monthly excess returns (including dividends) of ten stocks: A.G. Edwards, Citigroup, Morgan Stanley, Merrill Lynch, Dell, HP, IBM, Alcoa, Caterpillar, and P&G.
- Use the monthly series of 3-month Treasury bill rates of the secondary market as the risk-free interest rate to obtain simple excess returns.
- The sample span is from January 1990 to December 2006.



Asset Excess Returns



- 1 Review: Principal Component Analysis
- 2 **Kernel PCA**
- 3 Spectral Clustering

PCA Revisited

- Assume the data matrix \mathbf{X} is $N \times p$, we now allow $p \geq N$, which is usually called a “high-dimensional” problem.
- Set $r = \min\{p, N\}$.
- The SVD of \mathbf{X} has the form $\mathbf{X} = \mathbf{U}\mathbf{D}\mathbf{V}'$:
 - \mathbf{U} is $N \times r$, orthogonal.
 - \mathbf{V} is $p \times r$, orthogonal.
 - \mathbf{D} is $r \times r$, diagonal, with nonnegative diagonal entries.
- The principal components $z_1 = d_1 \mathbf{u}_1, \dots, z_r = d_r \mathbf{u}_r$ can be obtained without calculating the principal directions $\mathbf{v}_1, \dots, \mathbf{v}_r$.
 - Find the eigenvalues $\lambda_1 \geq \lambda_2 \geq \dots \geq \lambda_r$ and eigenvectors $\mathbf{u}_1, \dots, \mathbf{u}_r$ of the $N \times N$ matrix $\mathbf{X}\mathbf{X}'$.
 - The m -th principal component is given by $z_m = \sqrt{\lambda_m} \mathbf{u}_m$, for $1 \leq m \leq r$.
- The (i, j) -th entry of $\mathbf{X}\mathbf{X}'$ is $\langle \mathbf{x}_i, \mathbf{x}_j \rangle$.
- Suppose a new object \mathbf{x}_0 arrives, its principal components can also be calculated without knowing $\mathbf{v}_1, \dots, \mathbf{v}_r$.

$$z_{0m} := \mathbf{x}_0' \mathbf{v}_m = \mathbf{x}_0' \mathbf{X}' \mathbf{u}_m / d_m = \frac{1}{\sqrt{\lambda_m}} \sum_{i=1}^N u_{im} \langle \mathbf{x}_0, \mathbf{x}_i \rangle.$$

Derived Inputs and Kernel PCA

- Lift to a higher dimensional space using M -dimensional derived inputs $h(\mathbf{x}_i)$. Let $h(\mathbf{X}) = [h(\mathbf{x}_1), \dots, h(\mathbf{x}_N)]'$ be the $N \times M$ matrix with derived inputs.
- The principal components obtained using the derived inputs can be obtained from eigenvalues and eigenvectors of $h(\mathbf{X})[h(\mathbf{X})]'$.
- The (i, j) -th entry of $h(\mathbf{X})[h(\mathbf{X})]'$ is $\langle h(\mathbf{x}_i), h(\mathbf{x}_j) \rangle$.
- To calculate the principal components of a new object \mathbf{x}_0 , one only needs to calculate $\langle h(\mathbf{x}_0), h(\mathbf{x}_i) \rangle$.
- To calculate the principal components, all needed are the **inner products between (derived) input vectors**.
- Go one step further: replacing $\langle h(\mathbf{x}_i), h(\mathbf{x}_j) \rangle$ by $K(\mathbf{x}_i, \mathbf{x}_j)$, where $K(\cdot, \cdot)$ is a suitable kernel function.

Kernel PCA and Centering

- Let \mathbf{K} be the $N \times N$ matrix whose (i, j) -th entry is $K(\mathbf{x}_i, \mathbf{x}_j)$.
- Note that $h(\mathbf{x}_i)$ may not be centered, i.e. $\sum_{i=1}^N h(\mathbf{x}_i) \neq \mathbf{0}$.
- For the original PCA, easy to center the original inputs.
- For kernel PCA, how to center the infinite-dimensional derived inputs?
- Let $\mathbf{1} = \mathbf{1}_N$ be the N -dimensional vector with all entries equal to 1.

$$\left[\mathbf{I} - \frac{1}{N} \mathbf{1} \mathbf{1}' \right] \mathbf{x}^j = \mathbf{x}^j - \bar{x}_j \mathbf{1}.$$

- Let $\mathbf{J} = \mathbf{1} \mathbf{1}' / N$. Original PCA depends on the matrix

$$[(\mathbf{I} - \mathbf{J})\mathbf{X}][(\mathbf{I} - \mathbf{J})\mathbf{X}]' = (\mathbf{I} - \mathbf{J})(\mathbf{X}\mathbf{X}')(\mathbf{I} - \mathbf{J}).$$

- Similarly, in kernel PCA, use the matrix

$$\tilde{\mathbf{K}} = (\mathbf{I} - \mathbf{J})\mathbf{K}(\mathbf{I} - \mathbf{J}).$$

Kernel PCA Solution

- Suppose the rank of \tilde{K} is r .
- Let $\lambda_1 \geq \dots \geq \lambda_r$ and $\mathbf{u}_1, \dots, \mathbf{u}_r$ be the eigenvalues and eigenvectors of \tilde{K} .
- The m -th principal component is $\mathbf{z}_m = \sqrt{\lambda_m} \mathbf{u}_m$.
- Suppose a new object \mathbf{x}_0 arrives, its principal components are calculated as

$$z_{0m} = \frac{1}{\sqrt{\lambda_m}} \sum_{i=1}^N u_{im} \tilde{K}(\mathbf{x}_0, \mathbf{x}_i),$$

where

$$\tilde{K}(\mathbf{x}_0, \mathbf{x}_i) = \left[K(\mathbf{x}_0, \mathbf{x}_i) - \frac{1}{N} \sum_{j=1}^N K(\mathbf{x}_i, \mathbf{x}_j) - \frac{1}{N} \sum_{j=1}^N K(\mathbf{x}_0, \mathbf{x}_j) + \frac{1}{N^2} \mathbf{1}' \mathbf{K} \mathbf{1} \right].$$

- The principal components of a new object in matrix form is (arguably) simpler. It is left as a [homework problem](#).

Example

- Three clusters. The points are distributed uniformly on the circle, with radius 1, 2.8 and 5 in the three groups.
- Each coordinate of a point is added by a Gaussian noise with standard deviation 0.25.

```
n=150
```

```
p=2
```

```
set.seed(123)
```

```
X=matrix(rnorm(p*n*3),nrow=n*3)
```

```
X.norm=apply(X^2,MAR=1,FUN=sum)
```

```
X.norm=sqrt(X.norm)
```

```
X=X/X.norm
```

```
label=rep(c(1,2,3),each=n)
```

```
radius=rep(c(1,2.8,5),each=n)
```


```
color=rep(c("orange","blue","springgreen4"),each=n)
```

```
X=X*radius
```

```
X=X+.25*matrix(rnorm(p*n*3),nrow=n*3)
```


```
plot(X[,1],X[,2],col=color,xlab="1st Coordinate",ylab="2nd Coordinate")
```

Example: Data




Example: PCA

```
X.pca=svd(X)$u  
X.d=svd(X)$d  
plot(X.pca[,1]*X.d[1],X.pca[,2]*X.d[2],col=color,  
      xlab="1st Principal Component",ylab="2nd Principal Component")
```




Example: Kernel PCA 1

```
X.kpc1=kpca(X,kernel="rbfdot",kpar=list(sigma=.5),features=2)  
plot(rotated(X.kpc1),col=color,  
      xlab="1st Principal Component",ylab="2nd Principal Component")
```




Example: Kernel PCA 2

```
X.kpc2=kpca(X,kernel="rbfdot",kpar=list(sigma=.1),features=2)  
plot(rotated(X.kpc2),col=color,  
      xlab="1st Principal Component",ylab="2nd Principal Component")
```




Example: Kernel PCA 2 with a New Point

```
newx = predict(X.kpc2,t(c(0,0)))  
points(newx,pch=19,cex=2,col="red")
```




Example: Kernel PCA 3

```
X.kpc3=kpca(X,kernel="rbfdot",kpar=list(sigma=.01),features=2)  
plot(rotated(X.kpc3),col=color,  
      xlab="1st Principal Component",ylab="2nd Principal Component")
```




Excess Return: PCA

```
er10=read.table("04_er10.txt",header=T)
er10.mat=t(as.matrix(er10))
er10.pca=svd(er10.mat)
plot(er10.pca$u[,1],er10.pca$u[,2],xlab="1st Principal Component",ylab="2nd Principal Component")
text(er10.pca$u[,1],er10.pca$u[,2], rownames(er10.mat), cex=0.5, pos=2, col="blue")
```




Excess Return: Kernel PCA 1

```
kpc1=kpca(er10.mat,kernel="rbfdot",kpar=list(sigma=.4),features=2)
## quite stable over different sigma
plot(rotated(kpc1),xlab="1st Principal Component",ylab="2nd Principal Component")
text(rotated(kpc1), rownames(er10.mat), cex=0.5, pos=2, col="blue")
```



Excess Return: Kernel PCA 2

```
kpc2=kpca(er10.mat,kernel="polydot",kpar=list(degree=4,scale=1,offset=1),features=2)
## quite stable over different sigma
plot(rotated(kpc2),xlab="1st Principal Component",ylab="2nd Principal Component")
text(rotated(kpc2), rownames(er10.mat), cex=0.5, pos=4, col="blue")
```



- 1 Review: Principal Component Analysis
- 2 Kernel PCA
- 3 Spectral Clustering**

Similarity, Kernel, Similarity Graph

- Clustering: partition the observations $\{\mathbf{x}_1, \dots, \mathbf{x}_N\}$ into a few clusters.
- Traditional K-means does not work well when the clusters are non-convex.
- Spectral clustering is designed to better handle such a situation.
- Suppose there is a $N \times N$ matrix \mathbf{S} of pairwise similarities between all observation pairs.
 - $s_{ij} := \mathbf{S}[i, j]$ is the similarity between \mathbf{x}_i and \mathbf{x}_j .
 - E.g. $s_{ij} = 1/(d_{ij} + 1)$, where $d_{ij} = \|\mathbf{x}_i - \mathbf{x}_j\|^2$.
- Kernel method uses $s_{ij} = K(\mathbf{x}_i, \mathbf{x}_j)$, where $K(\cdot, \cdot)$ is a kernel function.
 - Radial basis kernel. $K(\mathbf{x}_i, \mathbf{x}_j) = \exp\{-\alpha\|\mathbf{x}_i - \mathbf{x}_j\|^2\}$.
- The observations can be represented in an **undirected similarity graph** $G = \langle V, E \rangle$.
 - Each observation is a vertex: \mathbf{x}_i corresponds to the vertex v_i .
 - The edge connecting \mathbf{x}_i and \mathbf{x}_j has the weight s_{ij} .
 - The graph G and the similarity matrix \mathbf{S} are equivalent.

Mutual K-Nearest-Neighbor Graph

- Suppose there is a $N \times N$ matrix of pairwise similarities s_{ij} between all observation pairs.
- Define the **Mutual K-nearest-neighbor graph** \mathbf{W} as follows.
 - \mathbf{W} is a $N \times N$ symmetric matrix. Denote its (i, j) -th entry by $w_{ij} := \mathbf{W}[i, j]$.
 - $w_{ii} = 0$.
 - For each pair (i, j) , if x_j is among the k -nearest neighbor of x_i , or vice versa, then define $w_{ij} = w_{ji} := s_{ij}$.
 - Otherwise, define $w_{ij} = w_{ji} = 0$.
 - The $N \times N$ matrix \mathbf{W} can be understood as a graph with N vertexes, where the edge between the i -th and j -th vertex is given by w_{ij} .
- Define $g_i = \sum_{j=1}^N w_{ij}$, and let \mathbf{G} be the diagonal matrix with diagonal elements $\{g_1, \dots, g_N\}$.
- The **Unnormalized graph Laplacian** is defined by $\mathbf{L} = \mathbf{G} - \mathbf{W}$.
 - Can also use the **normalized** version: $\tilde{\mathbf{L}} := \mathbf{I} - \mathbf{G}^{-1}\mathbf{W}$.

Spectral Clustering

- Spectral clustering finds the matrix $\mathbf{Z}_{N \times m}$, whose j -th column is the eigenvector corresponding to the $(N - j)$ -th largest eigenvalue of \mathbf{L} .
- In other words, the j -th column of \mathbf{Z} corresponds to the $(j + 1)$ -th smallest eigenvalue of \mathbf{L} .
- Now treat the i -th row as the new 'input' of the i -th point, and apply the standard K-means algorithm.
- How does it work? Let $\mathbf{f} \in \mathbb{R}^N$ be a unit vector,


$$\mathbf{f}'\mathbf{L}\mathbf{f} = \sum_{i=1}^N g_i f_i^2 - \sum_{i,j=1}^N f_i f_j w_{ij} = \frac{1}{2} \sum_{i,j=1}^N w_{ij} (f_i - f_j)^2.$$

- A pair (i, j) with a larger similarity w_{ij} will push the coordinates f_i and f_j closer.
- **Caution!** There is always an eigenvalue 0, with the constant eigenvector.
- If the graph \mathbf{W} is not connected, there will be more 0 eigenvalues.

Relationship with Kernel PCA


- The Kernel matrix K is the same as the similarity matrix S .
- The graph W is a localized version of S .
- Kernel PCA uses the centered version \tilde{K} of K .
- Kernel PCA finds eigenvector corresponding to the largest eigenvalues of \tilde{K} , which is equivalent to finding the eigenvectors corresponding to the smallest eigenvalues of $I - \tilde{K}$.
- Spectral clustering finds the eigenvectors corresponding to the smallest eigenvalues of $G - W$.

Example: Data



Example: K-Means

```
X.kmean=kmeans(X,centers=3)  
plot(X,col=X.kmean$cluster,xlab="1st Coordinate",ylab="2nd Coordinate")
```




Example: Spectral Clustering

```
##### Spectral clustering
knn.sym=function(X,k=5){
  n=dim(X)[1]
  X.dist=as.matrix(dist(X))
  NN.sym=array(0,c(n,n))
  for (i in 1:n){
    xorder=order(X.dist[i,])
    NN.sym[i,xorder[1:(k+1)]] = 1
  }
  NN.sym=NN.sym+t(NN.sym)
  NN.sym[NN.sym==2]=1
  return(NN.sym)
}

spec.clust=function(X,k=5,alpha=1){
  n=dim(X)[1]
  X.dist=as.matrix(dist(X))
  X.sim=exp(-alpha*X.dist^2) ##### Radial basis kernel with parameter alpha
  NN=knn.sym(X,k)
  W=X.sim*NN
  g=apply(W,MAR=1,FUN=sum)
  G=diag(g)
  L=G-W
  L.spd=eigen(L)
  return(L.spd)
}
```


Example: Eigenvalues

```
X.spc1=spec.clust(X,k=5,alpha=.1)  
n=dim(X)[1]  
plot(0:15,X.spc1$values[n-(0:15)],xlab="Number",ylab="Eigenvalue")  
abline(h=0,lty=2)
```




Example: 1st Smallest Eigenvector

```
plot(1:n,X.spc1$vertices[,n],xlab="Index",ylab="1st Smallest",col=color)
```




Example: 2nd Smallest Eigenvector

```
plot(1:n,X.spc1$vertices[,n-1],xlab="Index",ylab="2nd Smallest",col=color)
```




Example: 3rd Smallest Eigenvector

```
plot(1:n,X.spc1$vertices[,n-2],xlab="Index",ylab="3rd Smallest",col=color)
```



Example: 3rd vs 2nd

```
plot(X.spc1$variables[,n-1],X.spc1$variables[,n-2],  
     xlab="Second Smallest Eigenvector",ylab="Third Smallest Eigenvector",col=c
```



Example: K-Means Using Z

```
X.new=cbind(X.spc1$vectors[,n-1],X.spc1$vectors[,n-2])  
X.new.kmean=kmeans(X.new,centers=3)  
plot(X,col=X.new.kmean$cluster,xlab="1st Coordinate",ylab="2nd Coordinate")
```

