

Statistical Learning for Data Science

MSDS534

Lecture 07: Boosting

Department of Statistics & Biostatistics
Rutgers University

Acknowledgement

- Some of the figures in this presentation are taken from *An Introduction to Statistical Learning, with applications in R* (Springer, 2013) with permission from the authors: G. James, D. Witten, T. Hastie and R. Tibshirani.
- Some of the figures in this presentation are taken from *Elements of Statistical Learning* (Springer, 2009) with permission from the authors: T. Hastie, R. Tibshirani and J. Friedman.

Reading Assignments

- ESL: 10.1–10.13.1.

1 AdaBoost

2 Boosting Trees and Gradient Boosting

AdaBoost

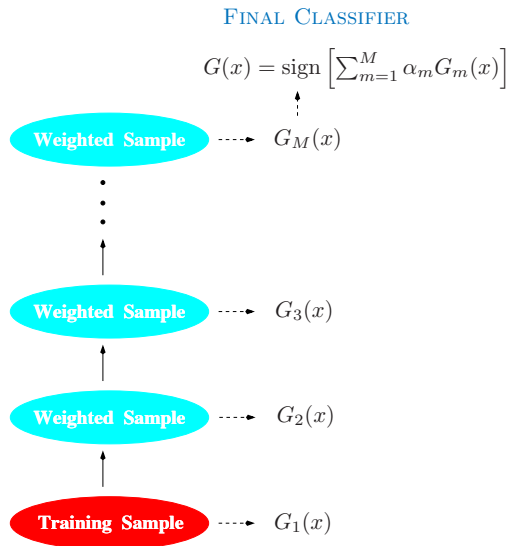


FIGURE 10.1. Schematic of AdaBoost. Classifiers are trained on weighted versions of the dataset, and then combined to produce a final prediction.

Algorithm 10.1 *AdaBoost.M1.*

1. Initialize the observation weights $w_i = 1/N$, $i = 1, 2, \dots, N$.
 2. For $m = 1$ to M :
 - (a) Fit a classifier $G_m(x)$ to the training data using weights w_i .
 - (b) Compute
$$\text{err}_m = \frac{\sum_{i=1}^N w_i I(y_i \neq G_m(x_i))}{\sum_{i=1}^N w_i}.$$
 - (c) Compute $\alpha_m = \log((1 - \text{err}_m)/\text{err}_m)$.
 - (d) Set $w_i \leftarrow w_i \cdot \exp[\alpha_m \cdot I(y_i \neq G_m(x_i))]$, $i = 1, 2, \dots, N$.
 3. Output $G(x) = \text{sign} \left[\sum_{m=1}^M \alpha_m G_m(x) \right]$.
-

Example

- Ten features X_1, \dots, X_{10} which are iid $N(0, 1)$.

-

$$Y = \begin{cases} 1 & \text{if } \sum_{j=1}^{10} X_j^2 \geq \chi_{10}^2(.5) \approx 9.34, \\ -1 & \text{otherwise.} \end{cases}$$

- Training sample size: 2,000. Test sample size: 10,000.
- Try the following:
 - Stump: a classification tree with two terminal nodes.
 - A 21-Node tree.
 - Boost the stump up to 400 iterations.

Example: Generate Data

```
library(tree)
library(gbm)

n.tr=2000
n.te=10000
p=10

set.seed(123)
X.tr=matrix(rnorm(n.tr*p),nrow=n.tr)
y=apply(X.tr^2,MAR=1,FUN="sum")
y=y>=9.34
y=as.factor(as.numeric(y))
ex1.tr=data.frame(X.tr,y)

X.te=matrix(rnorm(n.te*p),nrow=n.te)
y=apply(X.te^2,MAR=1,FUN="sum")
y=y>=9.34
y=as.factor(as.numeric(y))
ex1.te=data.frame(X.te,y)
```


Example: Trees


```
##### 21 Node Tree
ex1.tree=tree(y~.,ex1.tr)
summary(ex1.tree)
## Classification tree:
## tree(formula = y ~ ., data = ex1.tr)
## Number of terminal nodes: 21
## Residual mean deviance: 0.8235 = 1630 / 1979
## Misclassification error rate: 0.175 = 350 / 2000

ex1.tree.pred=predict(ex1.tree,ex1.te,type="class")
table(ex1.tree.pred,ex1.te$y)
test.error.tree=sum(ex1.tree.pred!=ex1.te$y)/n.te ##.2259

##### Stump
ex1.stump=prune.tree(ex1.tree,best=2,method="deviance")
## can replace "deviance" by "misclass"
## ex1.stump=prune.misclass(ex1.tree,best=2)
ex1.stump.pred=predict(ex1.stump,ex1.te,type="class")
table(ex1.stump.pred,ex1.te$y)
test.error.stump=sum(ex1.stump.pred!=ex1.te$y)/n.te ##.4649

##### Cross Validation
set.seed(3)
ex1.tree.cv =cv.tree(ex1.tree,FUN=prune.misclass)
ex1.tree.cv
plot(ex1.tree.cv$size,ex1.tree.cv$dev,type="b")
```

Example: Cross Validation



Example: Boost the stump


```
##### Boosting
ex1.tr$y=as.numeric(ex1.tr$y)-1
ex1.te$y=as.numeric(ex1.te$y)-1

ntree=400
## set.seed(0) ## Can set the seed if sampling is used
ex1.boost=gbm(y~.,data=ex1.tr,distribution="adaboost",n.trees=ntree,
              interaction.depth=1, shrinkage=1, bag.fraction=1)
summary(ex1.boost)
ex1.boost.pred=predict(ex1.boost,newdata=ex1.te[,1:p], n.trees=5)
table(ex1.boost.pred>=0,ex1.te$y)

test.error=rep(0,ntree)
for (i in 1:ntree){
  pp=predict(ex1.boost,newdata=ex1.te[,1:p], n.trees=i)
  pp=pp>=0
  pp=as.numeric(pp)
  test.error[i]=sum(pp!=ex1.te$y)/n.te
}

plot(1:ntree, test.error, type="l", col="orange",
     xlab="Boosting Iterations", ylab="Test Error" ,ylim=c(0,0.5))
abline(h=test.error.tree,lty=2,col="purple")
abline(h=test.error.stump,lty=2,col="gray")
legend("bottomleft", c("AdaBoost","Stump","21 Node Tree"),
      col=c("orange","gray","purple"), lty=c(1,2,2))
```

Example: Test Errors



Boosting Fits an Additive Model

- Suppose we want to learn a function $f(\mathbf{x})$ from the data.
 - $f(\cdot)$ can be a regression function.
 - For a binary classification, the rule can be given by the sign of $f(\cdot)$.
- Express $f(\cdot)$ through the basis expansion

$$f(\mathbf{x}) = \sum_{m=1}^M \beta_m b(\mathbf{x}; \gamma_m).$$

- β_m are expansion coefficients.
 - $b(\mathbf{x}; \gamma_m)$ is usually a simple function of \mathbf{x} , governed by a parameter vector γ_m .
- Examples of the function $b(\mathbf{x}; \gamma)$:
 - In neural networks, $b(\mathbf{x}, \gamma) = \sigma(\gamma_0 + \gamma'_1 \mathbf{x})$, where $\sigma(\cdot)$ is the sigmoid function.
 - In signal processing, $b(\mathbf{x}, \gamma)$ are wavelets, where γ parameterizes the location and scale shifts from a “mother” wavelet.
 - In nonparametric function estimation, $b(\mathbf{x}, \gamma)$ are splines, where γ depends on the knots.
 - For trees, γ parameterizes the split variables and split points at the internal nodes, and the predictions at the terminal nodes.

Forward Stagewise Additive Modeling

- Typically these models are fit by minimizing a loss function averaged over the training data

$$\min_{\{\beta_m, \gamma_m\}_1^M} \sum_{i=1}^N L \left(y_i, \sum_{m=1}^M \beta_m b(\mathbf{x}; \gamma_m) \right).$$

- For many loss functions $L(y, f(\mathbf{x}))$ and/or basis functions $b(\mathbf{x}; \gamma)$, this requires computationally intensive numerical optimization techniques.
- However, a simple alternative often can be found when it is feasible to rapidly solve the subproblem of fitting just a single basis function.
- **Forward stagewise modeling** approximates the solution by sequentially adding new basis functions to the expansion **WITHOUT** adjusting the parameters and coefficients of those that have already been added.

Forward Stagewise Additive Modeling

Algorithm 10.2 *Forward Stagewise Additive Modeling.*

1. Initialize $f_0(x) = 0$.

2. For $m = 1$ to M :

(a) Compute

$$(\beta_m, \gamma_m) = \arg \min_{\beta, \gamma} \sum_{i=1}^N L(y_i, f_{m-1}(x_i) + \beta b(x_i; \gamma)).$$

(b) Set $f_m(x) = f_{m-1}(x) + \beta_m b(x; \gamma_m)$.

Regression and Squared Loss

- Consider the squared-error loss

$$L(y, f(\mathbf{x})) = (y - f(\mathbf{x}))^2.$$

- For the m -th step,

$$\begin{aligned} L(y_i, f_{m-1}(\mathbf{x}_i) + \beta b(\mathbf{x}_i; \gamma_m)) &= (y_i - f_{m-1}(\mathbf{x}_i) - \beta b(\mathbf{x}_i; \gamma_m))^2 \\ &= (r_{im} - \beta b(\mathbf{x}_i; \gamma_m))^2. \end{aligned}$$

- $r_{im} = y_i - f_{m-1}(\mathbf{x}_i)$ is the residual of the current model on the i -th observation.
- Not only used for regression. Will also be used as an intermediate step in the gradient boosting algorithm, even for classification.

Algorithm 10.1 *AdaBoost.M1.*

1. Initialize the observation weights $w_i = 1/N$, $i = 1, 2, \dots, N$.
 2. For $m = 1$ to M :
 - (a) Fit a classifier $G_m(x)$ to the training data using weights w_i .
 - (b) Compute
$$\text{err}_m = \frac{\sum_{i=1}^N w_i I(y_i \neq G_m(x_i))}{\sum_{i=1}^N w_i}.$$
 - (c) Compute $\alpha_m = \log((1 - \text{err}_m)/\text{err}_m)$.
 - (d) Set $w_i \leftarrow w_i \cdot \exp[\alpha_m \cdot I(y_i \neq G_m(x_i))]$, $i = 1, 2, \dots, N$.
 3. Output $G(x) = \text{sign} \left[\sum_{m=1}^M \alpha_m G_m(x) \right]$.
-

Exponential Loss and AdaBoost

- The squared-error loss is generally not a good choice for classification.
- The AdaBoost.M1 (Algorithm 10.1) is equivalent to forward stagewise additive modeling (Algorithm 10.2) using the loss function

$$L(y, f(\mathbf{x})) = \exp(-yf(\mathbf{x})).$$

- $y \in \{-1, 1\}$.
 - If $y_i = 1$, and $f(\mathbf{x}_i) > 0$, then $\exp(-y_i f(\mathbf{x}_i))$ is small. And the larger $f(\mathbf{x}_i)$ is, the better.
- For the m -th step, we need to fit the classifier $G_m(\mathbf{x}) \in \{-1, 1\}$:

$$(\beta_m, G_m) = \arg \min_{\beta, G} \sum_{i=1}^N \exp \{-y_i [f_{m-1}(\mathbf{x}_i) + \beta G(\mathbf{x}_i)]\}.$$

- Can be re-written as

$$(\beta_m, G_m) = \arg \min_{\beta, G} \sum_{i=1}^N w_i^{(m)} \exp [-\beta y_i G(\mathbf{x}_i)].$$

- $w_i^{(m)} := \exp(-y_i f_{m-1}(\mathbf{x}_i))$ is a weight depending on neither β nor $G(\mathbf{x}_i)$.

Exponential Loss and AdaBoost

- The minimization over $G(\cdot)$ does **NOT** depend on β :

$$G_m = \arg \min_G \sum_{i=1}^N w_i^{(m)} I(y_i \neq G(\mathbf{x}_i)).$$

- We **haven't** specified what form $G(\cdot)$ takes. It can be, for example, a classification tree.
- Suppose the best G_m has been obtained, let err_m be the minimized weighted error rate

$$\text{err}_m = \frac{\sum_{i=1}^N w_i^{(m)} I(y_i \neq G_m(\mathbf{x}_i))}{\sum_{i=1}^N w_i^{(m)}}.$$

- It can be shown that the optimal β_m is given by

$$\beta_m = \frac{1}{2} \log \frac{1 - \text{err}_m}{\text{err}_m}.$$

- The classifier $f(\cdot)$ is updated as

$$f_m(\mathbf{x}) = f_{m-1}(\mathbf{x}) + \beta_m G_m(\mathbf{x}).$$

- The weights for the $(m+1)$ -th iteration becomes

$$w_i^{(m+1)} := w_i^{(m)} \cdot e^{-\beta_m y_i G_m(\mathbf{x}_i)} = w_i^{(m)} \cdot e^{\alpha_m I[y_i \neq G_m(\mathbf{x}_i)]} \cdot e^{-\beta_m}.$$


Exponential Loss and AdaBoost

- The training-set misclassification error decreases to zero at around 250 iterations (and remains there).
- The exponential loss keeps decreasing.
- The test-set misclassification error continues to improve after iteration 250.
- Adaboost is not optimizing training misclassification error; the exponential loss is more sensitive to changes in the estimated class probabilities.

```
train.error=rep(0,ntree)
exp.loss=rep(0,ntree)
for (i in 1:ntree){
  pp=predict(ex1.boost,newdata=ex1.tr[,1:p], n.trees=i)
  exp.loss[i]=mean(exp(-pp*(2*ex1.tr$y-1)))
  pp=pp>=0
  pp=as.numeric(pp)
  train.error[i]=sum(pp!=ex1.tr$y)/n.tr
}

plot(1:ntree, train.error, type="l", col="orange",
      xlab="Boosting Iterations", ylab="Test Error" ,ylim=c(0,1))
lines(1:ntree,exp.loss,type="l",col="purple")
legend("topright", c("Training Error","Exponential Loss"),
      col=c("orange","purple"), lty=c(1,1))
```

Example: Training Error and Exponential Loss



Why Exponential Loss?

- Under the conditional distribution of y given \mathbf{x}

$$f^*(\mathbf{x}) := \arg \min_{f(\mathbf{x})} \mathbb{E}_{y|\mathbf{x}}(e^{-yf(\mathbf{x})}) = \frac{1}{2} \log \frac{P(y = 1|\mathbf{x})}{P(y = -1|\mathbf{x})}.$$

- Or equivalently,

$$P(y = 1|\mathbf{x}) = \frac{1}{1 + e^{-2f^*(\mathbf{x})}}.$$

- The additive expansion produced by AdaBoost is estimating one-half of the log-odds of $P(y = 1|\mathbf{x})$.
- The **binomial negative log-likelihood**, or **deviance**, or **cross entropy** is

$$-\ell(y, f(\mathbf{x})) = \log \left(1 + e^{-2yf(\mathbf{x})} \right).$$

- It turns out

$$\arg \min_{f(\mathbf{x})} \mathbb{E}_{y|\mathbf{x}} [-\ell(y, f(\mathbf{x}))] = \arg \min_{f(\mathbf{x})} \mathbb{E}_{y|\mathbf{x}} (e^{-yf(\mathbf{x})}).$$

- The exponential loss and the binomial deviance lead to **DIFFERENT** results on finite data sets.

Loss Functions for Two-Class Classification




FIGURE 10.4. Loss functions for two-class classification. The response is $y = \pm 1$; the prediction is f , with class prediction $\text{sign}(f)$. The losses are misclassification: $I(\text{sign}(f) \neq y)$; exponential: $\exp(-yf)$; binomial deviance: $\log(1 + \exp(-2yf))$; squared error: $(y - f)^2$; and support vector: $(1 - yf)_+$ (see Section 12.3). Each function has been scaled so that it passes through the point (0, 1).

Compare Exponential Loss and Deviance Loss

- The classification rule is $G(\mathbf{x}) = \text{sign}[f(\mathbf{x})]$
 - Observations with positive margin $y_i f(\mathbf{x}_i) > 0$ are classified correctly.
 - Those with negative margin $y_i f(\mathbf{x}_i) < 0$ are misclassified.
 - The decision boundary is defined by $\{\mathbf{x} : f(\mathbf{x}) = 0\}$.
 - The goal of the classification algorithm is to produce positive margins as frequently as possible.
 - Any loss criterion should penalize negative margins more heavily than positive ones.
- Both the exponential and deviance loss can be viewed as monotone continuous approximations to misclassification loss.
 - The penalty associated with binomial deviance increases linearly for large increasingly negative margin.
 - The exponential criterion increases the influence of such observations exponentially.
- The exponential criterion concentrates much more influence on observations with large negative margins.
- The binomial deviance is more robust in noisy setting.
- The performance of AdaBoost has been empirically observed to dramatically degrade in such situations.

1 AdaBoost

2 Boosting Trees and Gradient Boosting

Boosting Trees

- Regression and classification trees partition the space of all joint predictor variable values into disjoint regions R_j , $j = 1, 2, \dots, J$, as represented by the terminal nodes of the tree.
- Let $\Theta = \{R_j, \gamma_j\}_1^J$ be the tree parameters.

$$T(\mathbf{x}; \Theta) = \sum_{j=1}^J \gamma_j I(\mathbf{x} \in R_j).$$

- J is usually treated as a meta-parameter. The regions R_j and the associated γ_j are obtained by minimizing the empirical loss

$$\hat{\Theta} = \arg \min_{\Theta} \sum_{j=1}^J \sum_{\mathbf{x}_i \in R_j} L(y_i, \gamma_j).$$

- This is a formidable combinatorial optimization problem, especially because of the regions.
- The boosted tree model is a sum of such trees

$$f(\mathbf{x}) = \sum_{m=1}^M T(\mathbf{x}; \Theta_m).$$

- Each tree $T(\mathbf{x}; \Theta_m)$ is usually simple, e.g. a stump.

Boosting Trees: AdaBoost

- The trees are induced in a forward stagewise manner (Algorithm 10.2).
- At the m -th iteration, the forward stagewise algorithm finds the m -th tree with parameters $\Theta_m = \{R_{jm}, \gamma_{jm}\}_1^{J_m}$, given the current $f_{m-1}(\mathbf{x})$.

$$\hat{\Theta}_m = \arg \min_{\Theta} \sum_{i=1}^N L[y_i, f_{m-1}(\mathbf{x}_i) + T(\mathbf{x}_i; \Theta_m)].$$

- AdaBoost finds the m -th tree that minimizes the weighted error rate

$$\sum_{i=1}^N w_i^{(m)} I(y_i \neq T(\mathbf{x}_i; \Theta_m)), \quad \text{where } w_i^{(m)} = \exp(-y_i f_{m-1}(\mathbf{x}_i)).$$

- In each region R_{jm} , γ_{jm} is either β_m or $-\beta_m$.
- This is a special kind of tree, which only takes two possible values on different regions. It's called a **scaled classification tree**.
- This is because we made the requirement that $G_m(\mathbf{x}) \in \{-1, 1\}$ at the beginning.

Boosting: General Case

- Without this restriction, the optimal constant γ_{jm} for the region R_{jm} is given by (typically easy to obtain)

$$\hat{\gamma}_{jm} = \arg \min_{\gamma_{jm}} \sum_{\mathbf{x}_i \in R_{jm}} L(y_i, f_{m-1}(\mathbf{x}_i) + \gamma_{jm}).$$

- Under the exponential loss,

$$\hat{\Theta}_m = \arg \min_{\Theta_m} \sum_{i=1}^N w_i^{(m)} \exp[-y_i T(\mathbf{x}_i; \Theta_m)].$$

- Given R_{jm} , the parameters γ_{jm} are then given by

$$\hat{\gamma}_{jm} = \frac{1}{2} \log \frac{\sum_{\mathbf{x}_i \in R_{jm}} w_i^{(m)} I(y_i = 1)}{\sum_{\mathbf{x}_i \in R_{jm}} w_i^{(m)} I(y_i = -1)}.$$

- Finding the region requires a specialized tree-growing algorithm. In practice, we prefer the approximation presented below that uses a weighted least squares regression tree.

Gradient Boosting

- An approximate algorithm for adding a tree with any **differentiable** loss criterion.
- The loss of using $f(x)$ to predict y on the training set is

$$L(f) = \sum_{i=1}^N L(y_i, f(x_i)).$$

- The goal is to minimize $L(f)$ with respect to f , where here $f(x)$ is constrained to be a sum of trees.
- Let $\mathbf{f} = [f(x_1), f(x_2), \dots, f(x_N)]'$. Then the problem becomes

$$\hat{\mathbf{f}} = \arg \min_{\mathbf{f}} L(\mathbf{f}).$$

- Using gradient descent, we try to reach or approximate $\hat{\mathbf{f}}$ by

$$\mathbf{f}_M = \sum_{m=0}^M \mathbf{h}_m, \quad \mathbf{h}_m \in \mathbb{R}^N.$$

- $\mathbf{f}_0 = \mathbf{h}_0$ is an initial guess.
- Each \mathbf{h}_m is induced based on the current \mathbf{f}_{m-1} .

Fit a Regression Tree to the Gradient

- Given the current \mathbf{f}_{m-1} , steepest descent (gradient descent) seeks to update \mathbf{f} along the direction \mathbf{g}_m , where

$$g_{im} = \left. \frac{\partial L(y_i, f_i)}{\partial f_i} \right|_{f_i = f_{m-1}(\mathbf{x}_i)}$$

- For optimization, can take $\mathbf{f}_m = \mathbf{f}_{m-1} - \rho_m \mathbf{g}_m$. The scalar ρ_m is determined by

$$\rho_m = \arg \min_{\rho} L(\mathbf{f}_{m-1} - \rho \mathbf{g}_m).$$

- However, we need a \mathbf{f}_m that is able to make predictions on new data.
- Induce a regression tree $T(\mathbf{x}; \Theta_m)$ by

$$\tilde{\Theta}_m = \arg \min_{\Theta} \sum_{i=1}^N [-g_{im} - T(\mathbf{x}_i; \Theta)]^2.$$

- After finding the regions from $\tilde{\Theta}_m$, the constants in each region are

$$\hat{\gamma}_{jm} = \arg \min_{\gamma_{jm}} \sum_{\mathbf{x}_i \in R_{jm}} L(y_i, f_{m-1}(\mathbf{x}_i) + \gamma_{jm}).$$

- Now update $\mathbf{f}_m(\mathbf{x}) = \mathbf{f}_{m-1}(\mathbf{x}) + \sum_{j=1}^{J_m} \gamma_{jm} I(\mathbf{x} \in R_{jm})$.

Gradient Tree Boosting Algorithm

Algorithm 10.3 *Gradient Tree Boosting Algorithm.*

1. Initialize $f_0(x) = \arg \min_{\gamma} \sum_{i=1}^N L(y_i, \gamma)$.

2. For $m = 1$ to M :

(a) For $i = 1, 2, \dots, N$ compute

$$r_{im} = - \left[\frac{\partial L(y_i, f(x_i))}{\partial f(x_i)} \right]_{f=f_{m-1}}.$$

(b) Fit a regression tree to the targets r_{im} giving terminal regions R_{jm} , $j = 1, 2, \dots, J_m$.

(c) For $j = 1, 2, \dots, J_m$ compute

$$\gamma_{jm} = \arg \min_{\gamma} \sum_{x_i \in R_{jm}} L(y_i, f_{m-1}(x_i) + \gamma).$$

(d) Update $f_m(x) = f_{m-1}(x) + \sum_{j=1}^{J_m} \gamma_{jm} I(x \in R_{jm})$.

3. Output $\hat{f}(x) = f_M(x)$.

Tree Size

- The simplest strategy is to restrict all trees to be of the same size J .
- No interaction effects of level greater than $J - 1$ are possible.
 - Setting $J = 2$ produces boosted models with only main effects.
 - With $J = 3$, two-variable interaction effects are also allowed.
 - The `gbm()` function has a parameter `interaction.depth` for the allowed level of interactions.
- Experience so far indicates that $4 \leq J \leq 8$ works well in the context of boosting, with results being fairly insensitive to particular choices in this range.

```
error.rate=function(m,newdata,ntree){  
  err=array(0,c(3,ntree))  
  rownames(err)=c("Mis","Exp","Dev")  
  for (i in 1:ntree){  
    p=dim(newdata)[2]-1  
    pp=predict(m,newdata=newdata[,1:p], n.trees=i)  
    err[2,i]=mean(exp(-pp*(2*newdata$y-1)))  
    err[3,i]=mean(log(1+exp(-2*pp*(2*newdata$y-1))))  
    pp=pp>=0  
    pp=as.numeric(pp)  
    err[1,i]=mean(pp!=newdata$y)  
  }  
  err  
}
```



Tree Size

```
##### Number of Nodes
m0=gbm(y~.,data=ex1.tr,distribution="adaboost",n.trees=ntree,
        interaction.depth=1, shrinkage=1, bag.fraction=1)
m1=gbm(y~.,data=ex1.tr,distribution="bernoulli",n.trees=ntree,
        interaction.depth=1, shrinkage=1, bag.fraction=1)
m2=gbm(y~.,data=ex1.tr,distribution="bernoulli",n.trees=ntree,
        interaction.depth=9, shrinkage=1, bag.fraction=1)
m3=gbm(y~.,data=ex1.tr,distribution="bernoulli",n.trees=ntree,
        interaction.depth=19, shrinkage=1, bag.fraction=1)

err0=error.rate(m0,newdata=ex1.te,ntree=ntree)
err1=error.rate(m1,newdata=ex1.te,ntree=ntree)
err2=error.rate(m2,newdata=ex1.te,ntree=ntree)
err3=error.rate(m3,newdata=ex1.te,ntree=ntree)

par(mar=c(4.5,4.5,.5,.4))
plot(1:ntree, err0[1,], type="l", col="orange", xlab="Boosting Iterations",
     ylab="Test Error" ,ylim=c(0,0.5))
lines(1:ntree,err1[1,],type="l",col="purple")
lines(1:ntree,err2[1,],type="l",col="gray")
lines(1:ntree,err3[1,],type="l",col="blue")
legend("topright", c("AdaBoost","Stump","10 Node", "20 Node"),
      col=c("orange","purple","gray","blue"), lty=rep(1,4))
```

Tree Size



Regularization: Shrinkage

- As with ridge regression and neural networks, shrinkage techniques can be employed as well.
- Scale the contribution of each tree by a factor $0 < \nu \leq 1$, i.e. line 2(d) of Algorithm 10.3 is replaced by

$$f_m(\mathbf{x}) = f_{m-1}(\mathbf{x}) + \nu \cdot \sum_{j=1}^J \gamma_{jm} I(\mathbf{x} \in R_{jm}).$$

- Smaller values of ν favor better test error, and require correspondingly larger values of M .
- The best strategy appears to be to set ν to be very small ($\nu < 0.1$) and then choose M by early stopping.
- The `gbm()` function has a parameter `shrinkage`, with default value 0.1.
- Example: the benefits of shrinkage are evident, especially when the binomial deviance is tracked.


Regularization: Shrinkage

```
##### Stump
##### Set interaction.depth=5 to experiment with the trees of 6 nodes
ntree=4000
m1=gbm(y~.,data=ex1.tr,distribution="bernoulli",n.trees=ntree,
        interaction.depth=1, shrinkage=1, bag.fraction=1)
err1=error.rate(m1,newdata=ex1.te,ntree=ntree)
m4=gbm(y~.,data=ex1.tr,distribution="bernoulli",n.trees=ntree,
        interaction.depth=1, shrinkage=.2, bag.fraction=1)
err4=error.rate(m4,newdata=ex1.te,ntree=ntree)

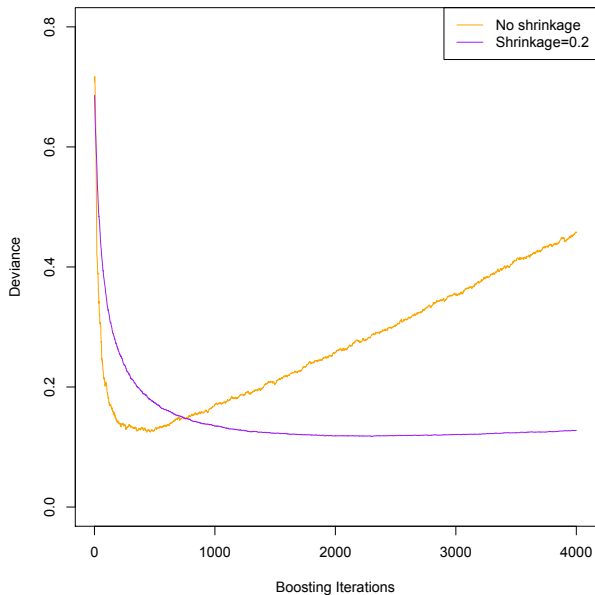
par(mar=c(4.5,4.5,.5,.4))
### Misclassification Error
plot(1:ntree, err1[1,], type="l", col="orange", xlab="Boosting Iterations",
     ylab="Test Error" ,ylim=c(0,0.5))
lines(1:ntree,err4[1,],type="l",col="purple")
legend("topright", c("No shrinkage","Shrinkage=0.2"),
     col=c("orange","purple"), lty=rep(1,2))

### Deviance
plot(1:ntree, err1[3,], type="l", col="orange", xlab="Boosting Iterations",
     ylab="Deviance" ,ylim=c(0,0.8))
lines(1:ntree,err4[3,],type="l",col="purple")
legend("topright", c("No shrinkage","Shrinkage=0.2"),
     col=c("orange","purple"), lty=rep(1,2))
```

Shrinkage: Stump



Shrinkage: Stump



Regularization: Subsampling

- With **stochastic gradient boosting**, at each iteration we sample a fraction η of the training observations (without replacement), and grow the next tree using that subsample.
- A typical value for η can be 0.5.
- For large N , η can be substantially smaller than 0.5.
- The `gbm()` function has a parameter `bag.fraction`, with default value 0.5
- Subsampling can reduce the computing time.
- In many cases it actually produces a more accurate model.

```


ntree=2000
m1=gbm(y~.,data=ex1.tr,distribution="bernoulli",n.trees=ntree,
        interaction.depth=1, shrinkage=1, bag.fraction=1)
err1=error.rate(m1,newdata=ex1.te,ntree=ntree)
m7=gbm(y~.,data=ex1.tr,distribution="bernoulli",n.trees=ntree,
        interaction.depth=1, shrinkage=0.2, bag.fraction=1)
err7=error.rate(m7,newdata=ex1.te,ntree=ntree)
m8=gbm(y~.,data=ex1.tr,distribution="bernoulli",n.trees=ntree,
        interaction.depth=1, shrinkage=1, bag.fraction=0.5)
err8=error.rate(m8,newdata=ex1.te,ntree=ntree)
m9=gbm(y~.,data=ex1.tr,distribution="bernoulli",n.trees=ntree,
        interaction.depth=1, shrinkage=0.2, bag.fraction=0.5)
err9=error.rate(m9,newdata=ex1.te,ntree=ntree)

### Misclassification Error
plot(1:ntree, err1[1,], type="l", col="orange", xlab="Boosting Iterations",
      ylab="Test Error" ,ylim=c(0,0.5))
lines(1:ntree,err7[1,],type="l",col="purple")
lines(1:ntree,err8[1,],type="l",col="gray")
lines(1:ntree,err9[1,],type="l",col="blue")
legend("topright",
      c("No shrinkage","Shrinkage=0.2","Sample=0.5", "Shrinkage=0.2 Sample=0.5"),
      col=c("orange","purple","gray","blue"), lty=rep(1,4))


### Deviance
plot(1:ntree, err1[3,], type="l", col="orange", xlab="Boosting Iterations",
      ylab="Deviance" ,ylim=c(0,0.8))
lines(1:ntree,err7[3,],type="l",col="purple")
lines(1:ntree,err8[3,],type="l",col="gray")
lines(1:ntree,err9[3,],type="l",col="blue")
legend("topright",
      c("No shrinkage", "Shrinkage=0.2", "Sample=0.5", "Shrinkage=0.2 Sample=0.5"),
      col=c("orange","purple","gray","blue"), lty=rep(1,4))

```


Shrinkage and Subsampling



Shrinkage and Subsampling



Relative Importance of Predictor Variables

- Often only a few of the predictors have substantial influence on the response; the vast majority are irrelevant and could just as well have not been included.
- For a single decision tree T with J nodes, Breiman et al. proposed


$$\mathcal{I}_\ell^2(T) = \sum_{t=1}^{J-1} \hat{\ell}_t^2 I[v(t) = \ell]$$

as a measure of relevance for each predictor variable X_ℓ .

- The sum is over the $J - 1$ internal nodes of the tree.
- At each internal node t , one of the input variables $X_{v(t)}$ is used to partition the region, which leads to the maximal improvement $\hat{\ell}_t^2$ in squared error loss.
- For an additive tree, the relative importance is defined as

$$\mathcal{I}_\ell^2 = \frac{1}{M} \sum_{m=1}^M \mathcal{I}_\ell^2(T).$$

Relative Importance



Example: Spam Data

Sources:

- (a) Creators: Mark Hopkins, Erik Reeber, George Forman, Jaap Suermondt
Hewlett-Packard Labs, 1501 Page Mill Rd., Palo Alto, CA 94304
- (b) Donor: George Forman (gforman at nospam hpl.hp.com) 650-857-7835
- (c) Generated: June-July 1999

- Number of Instances: 4601 (1813 Spam = 39.4%)
- The spam_ind.txt gives the label of training (0) and testing (0) data.
- The last column of denotes whether the e-mail was considered spam (1) or not (0).
- Number of Attributes (predictor variables): 57
- 48 continuous real [0,100] attributes of type word_freq_WORD = percentage of words in the e-mail that match WORD.
- 6 continuous real [0,100] attributes of type char_freq_CHAR = percentage of characters in the e-mail that match CHAR.
- 1 continuous real [1,...] attribute of type capital_run_length_average = average length of uninterrupted sequences of capital letters.
- 1 continuous integer [1,...] attribute of type capital_run_length_longest = length of longest uninterrupted sequence of capital letters.
- 1 continuous integer [1,...] attribute of type capital_run_length_total = total number of capital letters in the e-mail.

```
##### Example: Spam
spam=read.table("spam.txt",header=F)
colnames(spam)[58]="Y"
ind=scan("spam_ind.txt")
spam.tr=subset(spam,ind==0) ## size: 3065
dim(spam.tr)
spam.te=subset(spam,ind==1) ## size: 1536
dim(spam.te)

ntree=400
m1=gbm(Y~.,data=spam.tr,distribution="bernoulli",n.trees=ntree,
        interaction.depth=1,shrinkage=.1,bag.fraction=.5)
spam.pred=predict(m1,newdata=spam.te, n.trees=ntree)
table(spam.pred>=0,spam.te$Y) ## .054

ntree=400
m2=gbm(Y~.,data=spam.tr,distribution="bernoulli",n.trees=ntree,
        interaction.depth=3,shrinkage=.1,bag.fraction=.5)
spam.pred=predict(m2,newdata=spam.te, n.trees=ntree)
table(spam.pred>=0,spam.te$Y) ## .050

ntree=800
m3=gbm(Y~.,data=spam.tr,distribution="bernoulli",n.trees=ntree,
        interaction.depth=4,shrinkage=.1,bag.fraction=.5)
spam.pred=predict(m3,newdata=spam.te, n.trees=ntree)
table(spam.pred>=0,spam.te$Y) ## .042
summary(m3)
```

Spam: Relative Importance

